

Asymmetric Multiprocessing

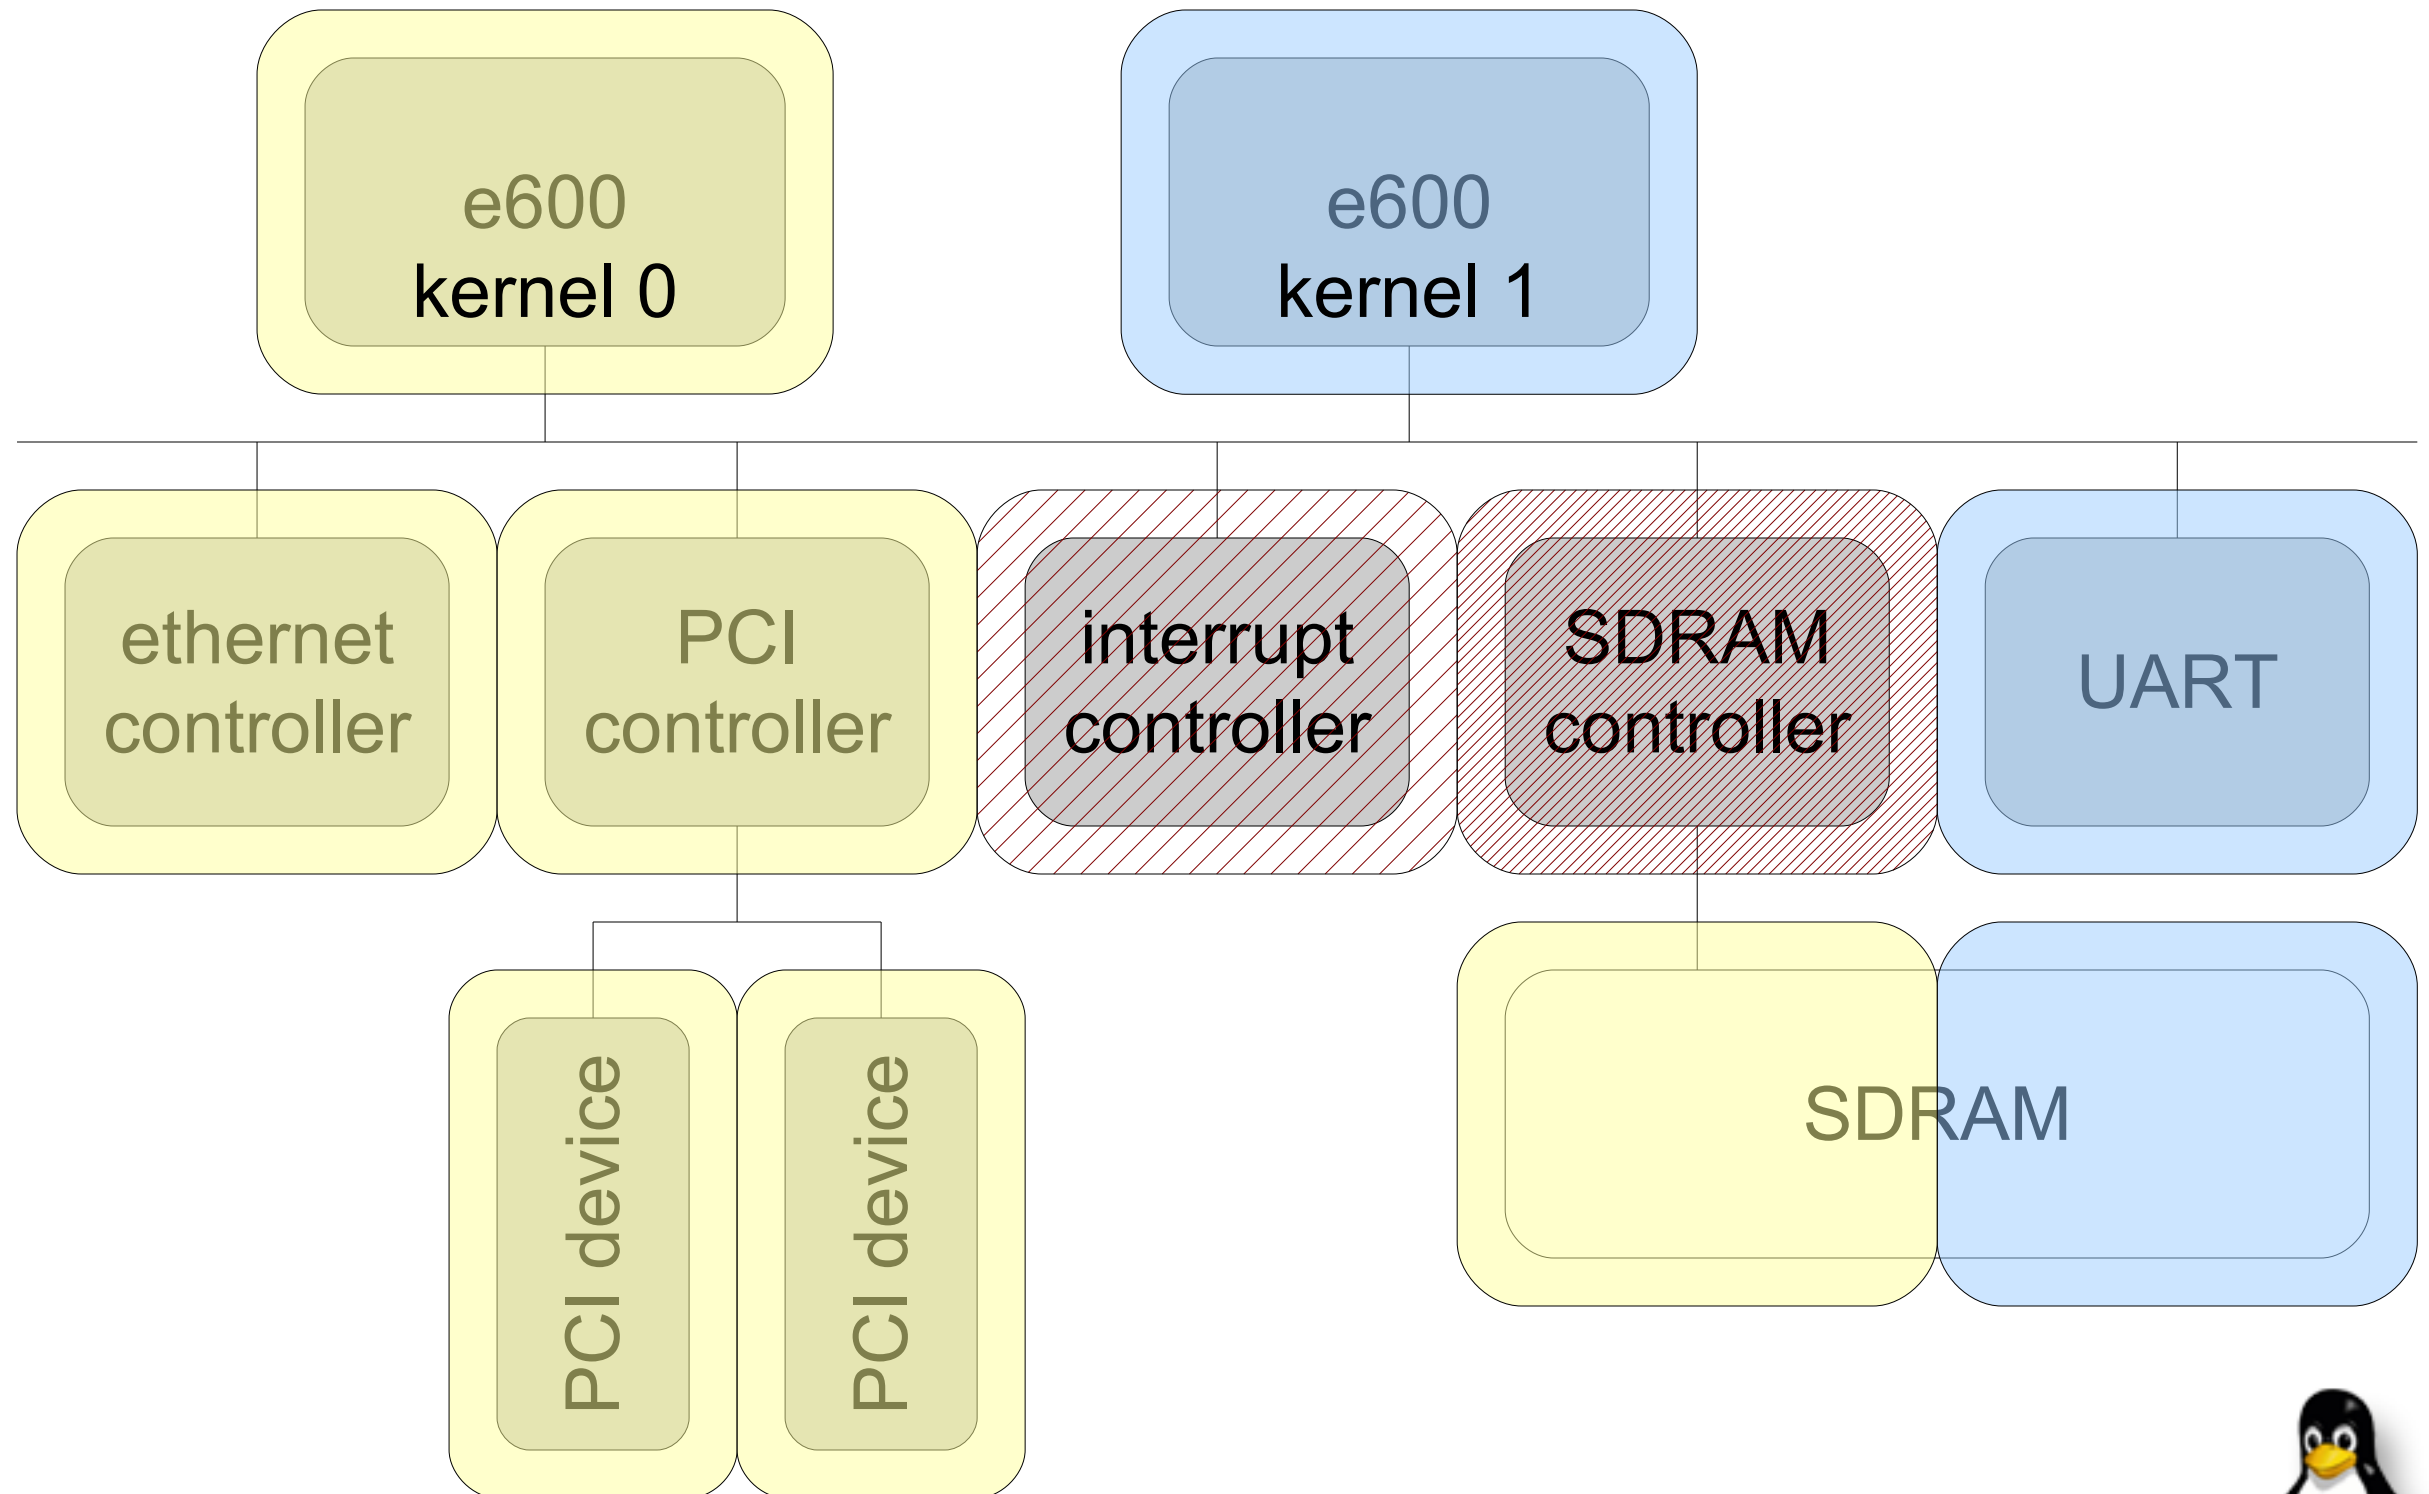
Linux Plumbers Conference

24 Sep 2009

Hollis Blanchard
IBM Linux Technology Center



Freescape 8641D SoC



Why do people want it?

- Multicore processors entering embedded space
- Linux + embedded OS
 - Enhance embedded OS capabilities with Linux UI, application availability, etc
 - Enhance Linux capabilities with dedicated-purpose software
 - Networking use case: control plane + data plane
 - One networking manufacturer claimed two orders of magnitude better throughput with data plane “custom executive”



What do you need?

- A way to load both kernels
 - U-boot – load both
 - CONFIG_OFFLINE – “unplug” resources from running Linux kernel
- Tell them what hardware (cores, memory, IO) they own
 - Device tree
- Modified kernels
 - E.g partial initialization of shared hardware
 - Non-zero RAM base
- Separate interrupt vectors for each core
- A way to share required hardware
- Inter-kernel communication



Shared Devices

- Devices initialized at boot
 - E.g. SDRAM controller
 - One kernel initializes, the other leaves it alone
 - Initializing kernel must have system-wide visibility to initialize correctly
 - Device tree
- Devices shared at runtime
 - Interrupt controller
 - But can't share interrupts (which core should be interrupted?)
 - Hardware must support concurrent accesses
 - e.g. write-to-clear registers for interrupt acknowledge



Inter-core Communication

- Could just use IO devices (e.g. 2 NICs + switch)
- Ideal: shared memory + interrupt
 - Starting to sound like virtio, right?
- Patches from Ira Snyder for virtio-over-pci
 - CompactPCI backplane



Related but Different: Embedded Virtualization

- Partitioning *with* isolation

- Good for reliability and debugging
- Even good for bringup: boot, crash, fix, repeat

- Issues

- Many embedded processors still lack hardware virtualization features
- Many embedded workloads are dominated by IO, which tends to suffer the greatest performance impact from virtualization



Summary

■ Pros

- Utilize multicore processors for power/space/BoM savings
- Use specialized kernels to solve different problems
- Doesn't require hardware virtualization support

■ Cons

- No isolation could cause very difficult debugging problems
- Possibly invasive kernel modifications required
- May need to duplicate IO devices
- May need a software proxying protocol to share IO devices

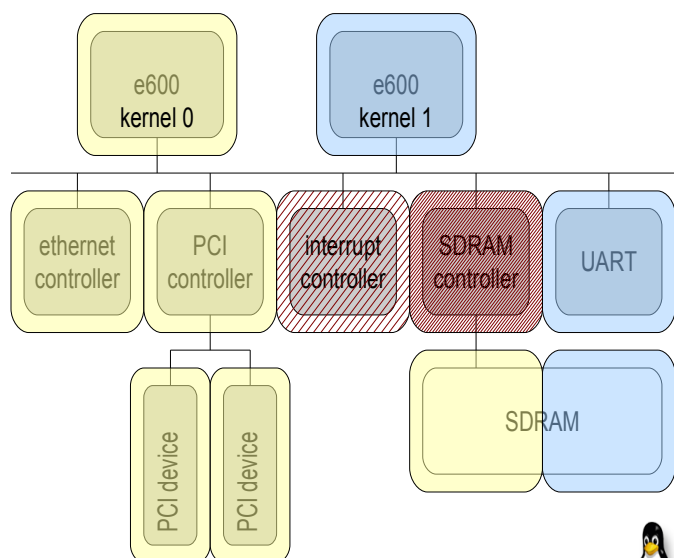


Asymmetric Multiprocessing
Linux Plumbers Conference
24 Sep 2009

Hollis Blanchard
IBM Linux Technology Center



Freescape 8641D SoC



Why do people want it?

- Multicore processors entering embedded space
- Linux + embedded OS
 - Enhance embedded OS capabilities with Linux UI, application availability, etc
 - Enhance Linux capabilities with dedicated-purpose software
 - Networking use case: control plane + data plane
 - One networking manufacturer claimed two orders of magnitude better throughput with data plane “custom executive”



What do you need?

- A way to load both kernels
 - U-boot – load both
 - CONFIG_OFFLINE – “unplug” resources from running Linux kernel
- Tell them what hardware (cores, memory, IO) they own
 - Device tree
- Modified kernels
 - E.g partial initialization of shared hardware
 - Non-zero RAM base
- Separate interrupt vectors for each core
- A way to share required hardware
- Inter-kernel communication



Shared Devices

- Devices initialized at boot
 - E.g. SDRAM controller
 - One kernel initializes, the other leaves it alone
 - Initializing kernel must have system-wide visibility to initialize correctly
 - Device tree
- Devices shared at runtime
 - Interrupt controller
 - But can't share interrupts (which core should be interrupted?)
 - Hardware must support concurrent accesses
 - e.g. write-to-clear registers for interrupt acknowledge



Inter-core Communication

- Could just use IO devices (e.g. 2 NICs + switch)
- Ideal: shared memory + interrupt
 - Starting to sound like virtio, right?
- Patches from Ira Snyder for virtio-over-pci
 - CompactPCI backplane



Related but Different: Embedded Virtualization

- Partitioning *with* isolation
 - Good for reliability and debugging
 - Even good for bringup: boot, crash, fix, repeat
- Issues
 - Many embedded processors still lack hardware virtualization features
 - Many embedded workloads are dominated by IO, which tends to suffer the greatest performance impact from virtualization



Summary

■ Pros

- Utilize multicore processors for power/space/BoM savings
- Use specialized kernels to solve different problems
- Doesn't require hardware virtualization support

■ Cons

- No isolation could cause very difficult debugging problems
- Possibly invasive kernel modifications required
- May need to duplicate IO devices
- May need a software proxying protocol to share IO devices

