


Why USB 3.0?

- 480Mb/s is too slow
- USB 2.0 sucks power
 - Inefficient host controller design
 - Polling and broadcast messages
 - Many devices don't support auto-suspend
- Start of frames (SOFs) sent with one active device



Why is USB 3.0 interesting?

- Backwards compatible
- Faster speed (5Gbps) with room to grow

4	<i>wSpeedsSupported</i>	2	Bitmap	Bitmap encoding of the speed supported by this device when operating in SuperSpeed mode. <table><tr><td><u>Bit</u></td><td><u>Encoding</u></td></tr><tr><td>0</td><td>If this bit is set, then the device supports operation at low-Speed USB.</td></tr><tr><td>1</td><td>If this bit is set, then the device supports operation at full-Speed USB.</td></tr><tr><td>2</td><td>If this bit is set, then the device supports operation at high-Speed USB.</td></tr><tr><td>3</td><td>If this bit is set, then the device supports operation at 5 Gbps.</td></tr><tr><td>15:4</td><td>Reserved. Shall be set to zero.</td></tr></table>	<u>Bit</u>	<u>Encoding</u>	0	If this bit is set, then the device supports operation at low-Speed USB.	1	If this bit is set, then the device supports operation at full-Speed USB.	2	If this bit is set, then the device supports operation at high-Speed USB.	3	If this bit is set, then the device supports operation at 5 Gbps.	15:4	Reserved. Shall be set to zero.
<u>Bit</u>	<u>Encoding</u>															
0	If this bit is set, then the device supports operation at low-Speed USB.															
1	If this bit is set, then the device supports operation at full-Speed USB.															
2	If this bit is set, then the device supports operation at high-Speed USB.															
3	If this bit is set, then the device supports operation at 5 Gbps.															
15:4	Reserved. Shall be set to zero.															

- Bulk “streams” allow SCSI command queuing

Why is USB 3.0 interesting?

- Better power management
 - device notifications (no more polling)
 - unicast packets (not broadcast)
 - link power management
 - function power management
 - host controller schedule in HW, not system memory



USB 3.0 Implications

- 6 wires added for USB 3.0
- USB 2.0 devices use separate wires
 - Same PM/auto-suspend problems as before
- New host controller (xHCI), new host controller driver
 - scheduler in hardware, xHCI driver needs hooks for device changes




Figure 5-15. Illustration of a USB 3.0 Cable Cross-Section

USB 3.0 host-side cable (standard A)






Figure 5-5. Illustration of Color Coding Recommendation for USB 3.0 Standard

USB 3.0 device side cable (standard-B and mini-B)



State of xHCI/USB 3.0 in Linux

- Supported in 2.6.31:
 - device enumeration
 - bulk and control TX
 - all device speeds (LS/FS/HS/SS)
 - stalls
 - cancellation
- Ready for 2.6.32:
 - interrupt TX
 - devices under 2.0 hubs
 - babbles




YES

SOON



MAYBE



NOT YET



SOON


xHCI driver future changes

- setting alternate interfaces
- isochronous TX
- non-standard polling rates
- resetting devices
- little endian support
- USB 3.0 bulk streams
- USB autosuspend
- xHCI PCI device suspend
- virtualization



Kernel Changes separate from xHCI

- New USB device class drivers
- USB 3.0 hub support
- USB 3.0 Function PM
- USB 3.0 Link PM
- Can you help with these?



Current USB power management

- Automatically suspend the whole device
- Userspace must enable auto-suspend
- Drivers must support auto-suspend
- USB core keeps track of idleness
- Devices have to not break!



USB 3.0 function PM

- USB 2.0 has device suspend
 - suspend whole USB device
- USB 3.0 also has device suspend, but it adds function suspend
 - suspend a set of related interfaces on a device
 - use IAD to find related interfaces



OS changes for USB 3.0 function PM

- USB core needs to handle function PM
- Track when an interface is claimed or busy
- Use Interface Association Descriptor (IAD)
- Send function suspend when interfaces are idle
- Handle Function Wake Device Notifications
- Putting all functions into suspend does not put the **device** into suspend; still need to send device suspend request


USB 3.0 Link PM

- USB 3.0 traffic is unicast
- Each idle link can be put into lower-power states (U0, U1, U2)
- Each link state has an exit latency
- Sort of like CPU C-states
- Each link partner can ask to go into a lower link state
- Highest link state is propagated up




OS changes for USB 3.0 Link PM

- Hardware does most of the work
- Software needs to set backup policy
- Need to set U1/U2 timeouts for each hub port
- Need some “wiggle room” in timeouts - maybe 5 to 10 times max exit latency?



OS changes for USB 3.0 Link PM

- Decide if it's worth it to enable U1/U2 for a device
- Is a periodic device too deep in the device tree?
- Are the hubs too slow?



OS changes for USB 3.0 Link PM

- Most of the work in USB core
- xHCI will trap roothub timeouts
- xHCI needs to set the maximum propagation delay for each device

USB 3.0 hubs

- Changes need to be made to khubd
 - new device descriptor
 - new class-specific requests
 - different port status bits
 - no transaction translators
 - hot reset vs. warm reset

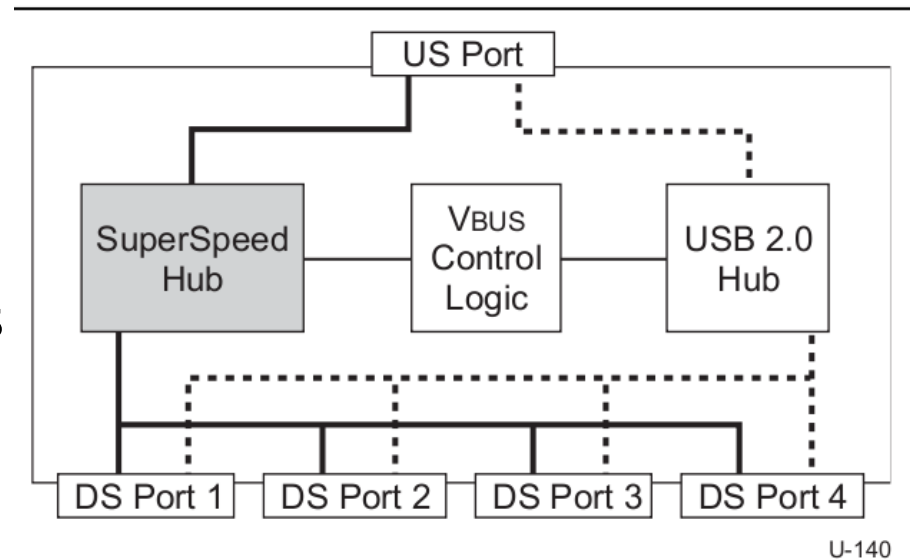


Figure 10-1. Hub Architecture

USB 3.0 Bulk “streams”

- Some USB 3.0 bulk endpoints support multiple “streams”
- Packets are tagged with a stream ID
- Device is notified when a stream has new data
- Device can start and stop any stream it wants to

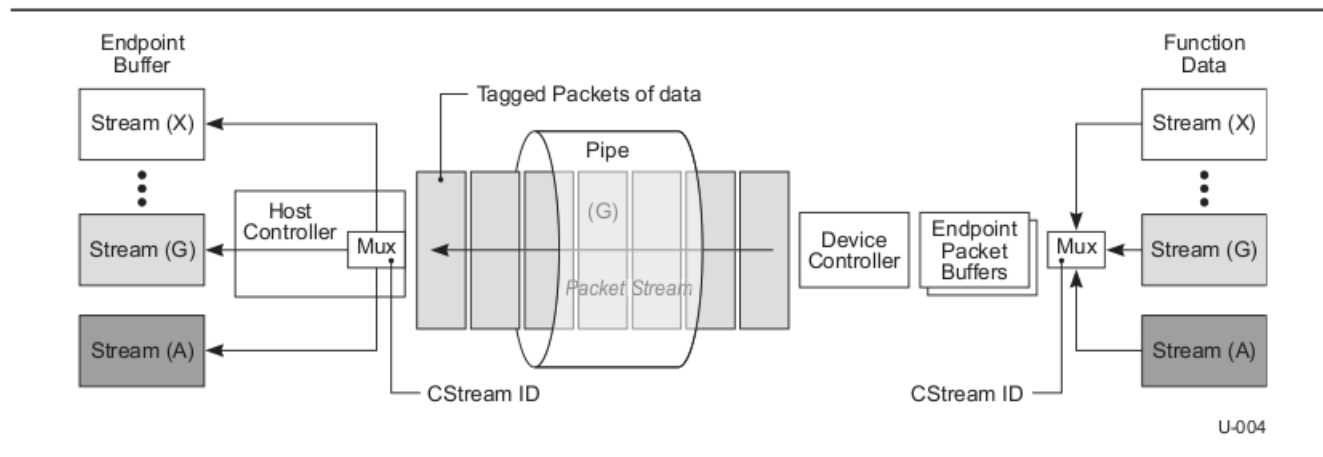


Figure 4-3. USB SuperSpeed IN Stream Example

USB 3.0 Bulk “streams”

- Allows each SCSI command to be tagged with a stream ID
- MSC device decides which command to start
- Spinning disks can sort commands
- Flash & SSDs can start prefetching sooner




Figure 4-3. USB SuperSpeed IN Stream Example

USB 3.0 storage devices

- Some will be legacy (BOT)
- USB Attached SCSI Protocol (UASP)
- Can be a USB 2.0 or USB 3.0 device
- Uses USB 3.0 bulk streams to queue multiple SCSI commands to device
- New USB class driver
- xHCI needs to support bulk streams



USB 3.0 webcams

- Point Grey webcam announced at IDF
- uncompressed 1080p video
- Will V4L layer handle this?
- Some USB video drivers have assumptions based on speed
 - e.g. driver picks a different polling interval based on FS or HS



Kernel/Userspace Interface changes for USB 3.0

- usbfs and libusb need to become aware of USB 3.0 stream IDs.
- Is it fast enough? Do we need a scatter-gather interface?
- USBMon needs to understand scatter gather lists and stream IDs.



Userspace changes for USB 3.0

- New UASP class with SCSI command queuing should have little impact on userspace
- How will applications like cheese handle faster USB webcams?
- Is HAL ready for USB 3.0?

How can I help?

- Areas you can help in:
 - New USB device class drivers
 - Readyng old class drivers for USB 3.0 devices
 - USB 3.0 hub support
 - USB 3.0 Link PM
 - USB 3.0 Function PM
- Patches and discussion on the Linux USB mailing list:
 - linux-usb@vger.kernel.org
 - <http://www.linux-usb.org/mailing.html>
- xHCI git tree on kernel.org

Questions?

Sarah Sharp

sarah.a.sharp@linux.intel.com

twitter: @sarahsharp

Creative Commons Attributions

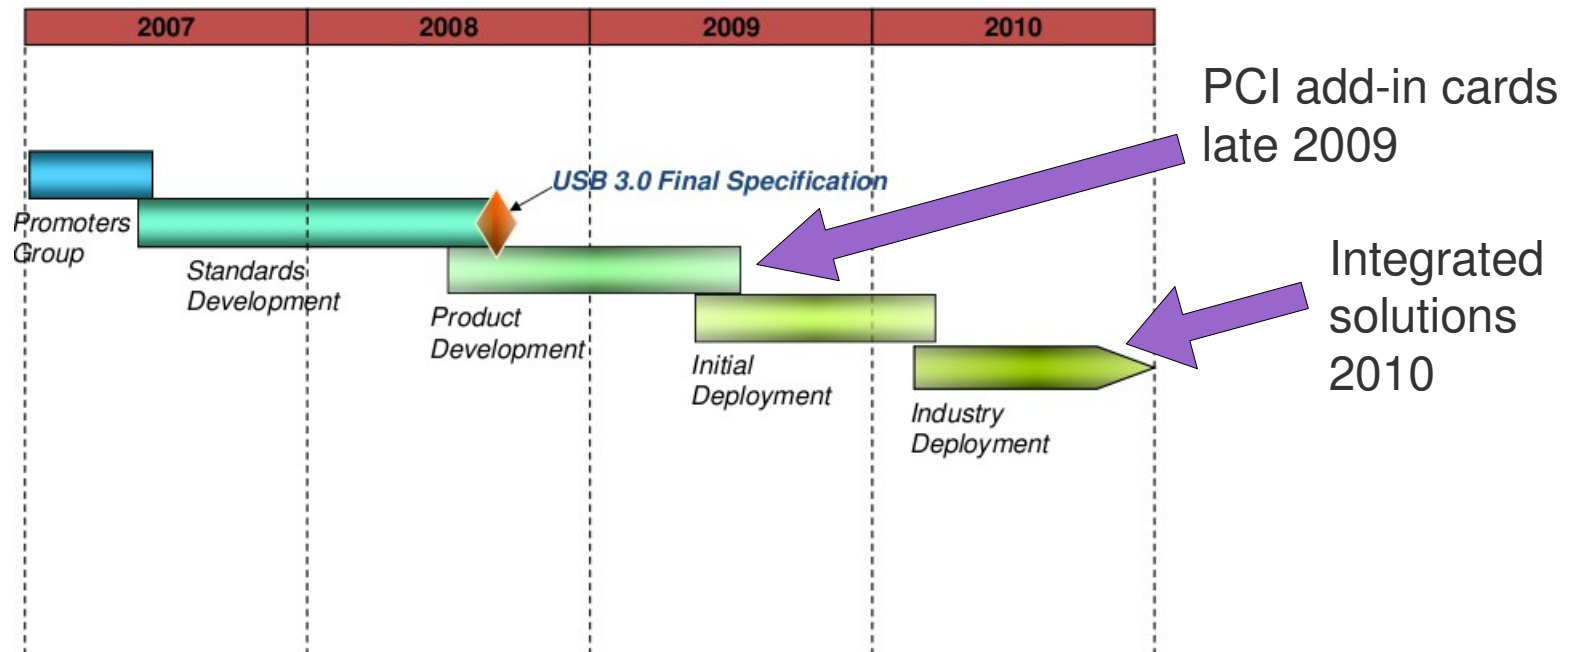
- light bulb http://www.flickr.com/photos/mr_beaver/3486761520/
- “W” thumb drive <http://www.flickr.com/photos/ivanwalsh/3657331200/>
- keyboard & mice <http://www.flickr.com/photos/m0php/3862857014/>
- webcam <http://www.flickr.com/photos/mrtea/772346725/>
- blue hub <http://www.flickr.com/photos/jeanbaptistem/3486039048/>
- work in progress <http://www.flickr.com/photos/hellochris/2801931497/>
- host ports <http://www.flickr.com/photos/kikus/3732845777/>
- purple thumb drive <http://www.flickr.com/photos/caroslines/2046327031/>
- USB to SATA <http://www.flickr.com/photos/cavemonkey50/427366996/>
- webcam under linux <http://www.flickr.com/photos/phylevn/2948896990/>
- plumbers nightmare <http://www.flickr.com/photos/ejbsf/3413576188/>
- printer <http://www.flickr.com/photos/davesag/192584714/>
- are we there yet? <http://www.flickr.com/photos/caseya/372922053/>
- logitech mouse <http://www.flickr.com/photos/blogitech/2883630458/>

Other photos

- USB 3.0 devices at IDF from engadget and reghardware

When will USB 3.0 devices appear?


- Jeff Ravencraft's (USB-IF Pres.) estimated timeline:



- NEC announced certified discrete host controller

Upsides of USB 3.0: No more polling

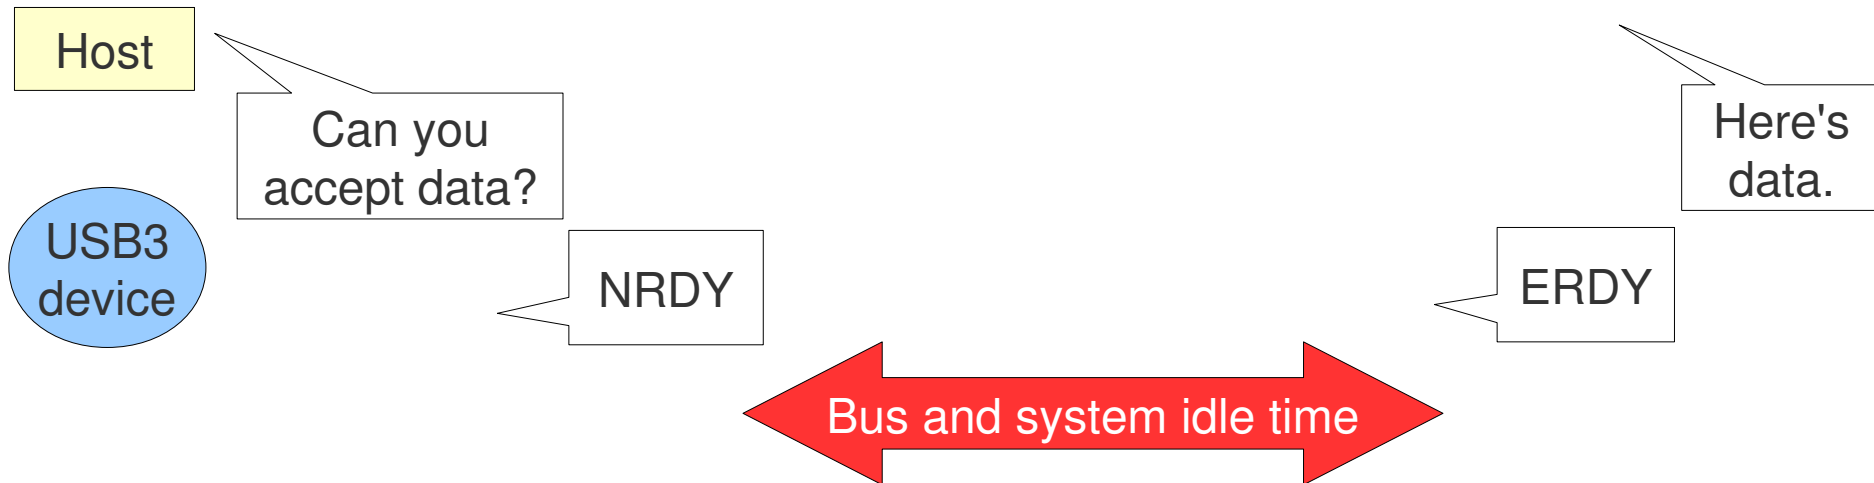
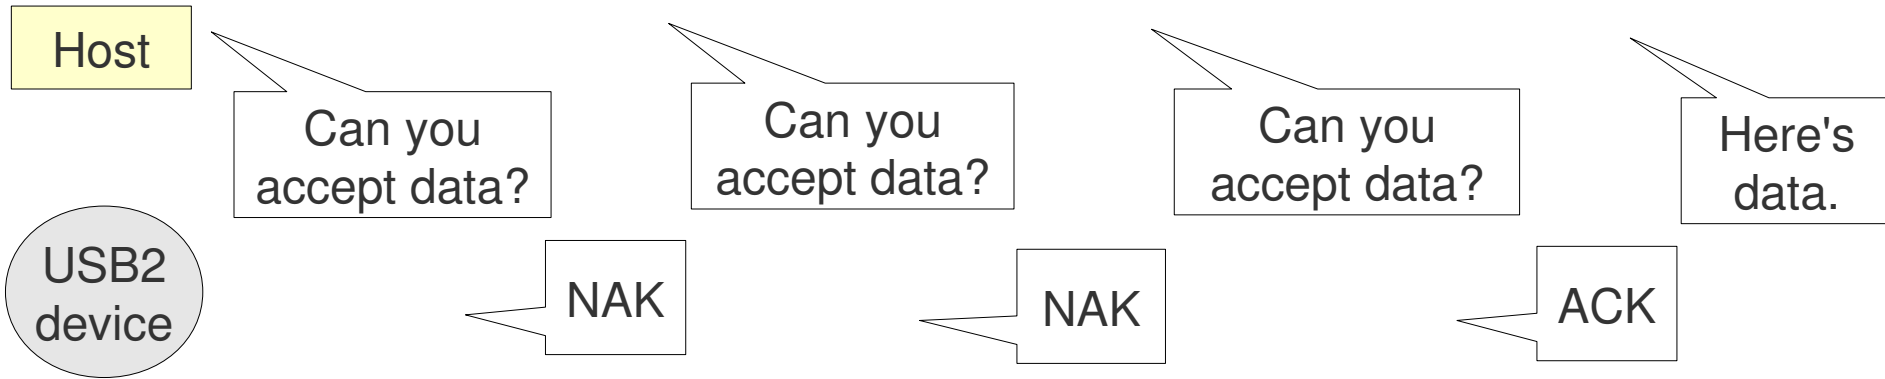
- High, full, and low speed devices can NAK an OUT transfer if they aren't ready to process the data.
- Leads to a lot of bus activity.
- USB 3.0 devices can say they aren't ready for data yet (NRDY)
- When they are ready, they asynchronously notify the host (ERDY)



Are we there yet?



USB 2.0 polling vs. USB 3.0 NRDY/ERDY



Implications of USB 3.0 NRDY/ERDY

- EHCI sets NAK count to 4
 - host controller gives up after 4 NAKs
 - max wait time of 4ms for FS/LS device response
- xHCI has no timeout on NRDY'ed transfers
 - Could be on the order of seconds?
- Implication: Userspace shouldn't block on USB transactions
 - X polling /dev/eventN for mouse movement - should be fine since it uses fnotify (and no one will make a USB3 mouse)
 - What about HAL polling?

USB 3.0 Link Power Management

- Routed packets means some bus links will be idle
- Two new link power management states
- Deeper power savings and higher exit latencies

Table C-1. Link States and Characteristics Summary

Link State	Description	Characteristics	State Transition Initiator	Device Clock Gen On/Off	Typical Exit Latency Range
U0	Link active	Link operational state	N/A	On	N/A
U1	Link idle – fast exit	Rx and Tx circuitry quiesced	Hardware ¹	On or Off	µs
U2	Link idle – slower exit	Clock generation circuitry may additionally be quiesced	Hardware ¹	On or Off ²	µs – ms
U3	Link suspend	Interface (e.g., Physical Layer) power may be removed	Entry: Software only Exit: Hardware or Software	Off	ms

Notes:

1. It is possible, under system test conditions, to instrument software initiated U1 and U2 state transitions.
2. From a power efficiency perspective it is desirable for devices to turn off their clock generation circuitry (e.g., their PLL) during the U2 link state.