

# Threaded Network Interrupts

Steven Rostedt

[srostedt@redhat.com](mailto:srostedt@redhat.com)

[<rostedt@goodmis.org>](mailto:rostedt@goodmis.org)

<http://people.redhat.com/srostedt>  
network-thread-irq.odp

# Disclaimer

I'm not a network guru

I had to fight for the hardware I have

# Threaded Interrupts

- Similar to NAPI code
- Interrupt comes in, disable the device
- Thread is awoken
- Thread polls on device, processing packets
- Finished and enables device

# Benchmarking


- On my 1Gb NICs, all kernels had same result in netperf
  - Saturated network
- 10 Gig NIC
  - Could not get working with the threaded infrastructure
    - I'm not a network device expert
  - Decided to just give up and test with the RT patch
    - This is far from optimizing the design


# Benchmark Machine

- Two Nehalem machines
- 2 x 4 core Intel Xeons
- Chelsio Communications Inc T310 10GbE Single Port Adapter


# Benchmark


- Ran netperf -H host
- Red Hat Enterprise Linux 5.4
- Full Preempt Real Time Patch (PREEMPT\_RT)
  - 2.6.31-rt10
- Hard and Soft interrupts as threads
  - RT patch without PREEMPT\_RT enabled
- Soft interrupts as threads
  - RT patch with only PREEMPT\_SOFTIRQ set
- Vanilla Linus Kernel
  - 2.6.31 downloaded from [kernel.org](http://kernel.org)


































# Cons

- Horrible latency on server config (no kernel preemption)
- Redesign of network infrastructure

# Pros

- Can schedule in handler
- Easier locking algorithms
- Prioritize interrupts
- Smoother desktop experience

# Discussion

# Threaded Network Interrupts

Steven Rostedt  
[srostedt@redhat.com](mailto:srostedt@redhat.com)  
<[rostedt@goodmis.org](mailto:rostedt@goodmis.org)>  
<http://people.redhat.com/srostedt>  
network-thread-irq.odp

# Disclaimer

I'm not a network guru

I had to fight for the hardware I have

# Threaded Interrupts

- Similar to NAPI code
- Interrupt comes in, disable the device
- Thread is awoken
- Thread polls on device, processing packets
- Finished and enables device

# Benchmarking

- On my 1Gb NICs, all kernels had same result in netperf
  - Saturated network
- 10 Gig NIC
  - Could not get working with the threaded infrastructure
    - I'm not a network device expert
  - Decided to just give up and test with the RT patch
    - This is far from optimizing the design





# Benchmark Machine


- Two Nehalem machines
- 2 x 4 core Intel Xeons
- Chelsio Communications Inc T310 10GbE Single Port Adapter


# Benchmark


- Ran netperf -H host
- Red Hat Enterprise Linux 5.4
- Full Preempt Real Time Patch (PREEMPT\_RT)
  - 2.6.31-rt10
- Hard and Soft interrupts as threads
  - RT patch without PREEMPT\_RT enabled
- Soft interrupts as threads
  - RT patch with only PREEMPT\_SOFTIRQ set
- Vanilla Linus Kernel
  - 2.6.31 downloaded from kernel.org































## Cons

- Horrible latency on server config (no kernel preemption)
- Redesign of network infrastructure

## Pros

- Can schedule in handler
- Easier locking algorithms
- Prioritize interrupts
- Smoother desktop experience

# Discussion