



# Real-Time Micro-Benchmark for Real-Time Linux

Vernon Mauery

---

IBM Linux Technology Center  
Real-Time Linux



## Disclaimer

This is not my code, but I know who wrote it

Mike Fulton <fultonm@ca.ibm.com>

Ajith Ramanath <ajith.ramanath@in.ibm.com>

Radhakrishnan Thangamuthu <rthangam@in.ibm.com>

Recently released into the wild

<http://sourceforge.net/projects/rtmicrobench/>

for those who just can't wait:

svn co <https://rtmicrobench.svn.sourceforge.net/svnroot/rtmicrobench/rtmb>



# Benchmarking

Lies, damn lies, and benchmarks

Fabricated workloads that give a skewed picture

Can offer a small window into specific workloads

Pick your favorite benchmark to tell you the lies that sooth you the most



# Yet Another Benchmark?

No ideal RT benchmark

- Cross platform

- Cross language

- Testing RT APIs (with micro benchmarks)

Customer Java acceptance test

- Does it fail in C too?

- A small pile of disembodied tests

- Still testing against hard-coded standards

Starting from the bottom up

- Use an automatic calibration

- Use a test harness

- Use proper statistical models



## How does it work?

Establish the clock resolution

For each remaining test

- For iteration count

  - Get current time

  - Perform work (test some API or other small task)

  - Get time

- Did the test run for longer than the clock resolution?

  - If not, increase work size

- Is determinism within user constraints? (soft, hard, guaranteed)

  - re-run with larger workload (longer time) or quit

- Give the test a latency score based on how deterministic it is

- Give a throughput score at that level of determinism




# Determinism Scores

Soft RT

30% deviation


99<sup>th</sup> percentile



Hard RT

20% deviation


99.9<sup>th</sup> percentile



Guaranteed RT

10% deviation

99.999<sup>th</sup> percentile





## What does it test?

	C	C++	Java
Bytecode Consistency			X
Clock	X	X	X
Compilation			X
Disk I/O read	X	X	X
Disk I/O write	X	X	X
Event dispatch	X	X	X
Float rate	X	X	X
Garbage generation			X
Int rate	X	X	X
Lock consistency	X	X	X
Mcache perf	X	X	X
Multi CPU		X	X
Multi thread	X	X	X
Net I/O	X	X	X
NHRT support			X
Periodic event	X	X	X
Scoped memory			X
Thread priority	X	X	X
Time accuracy	X	X	X
Timer	X	X	X





# Tests

## Bytecode Consistency (Java)

Verifies that each bytecode instructions perform correctly

## Clock (All)

Determines the granularity of the clock

## Compilation (Java)

Verifies that the bytecode compiler does not interfere with the performance of the application

## Disk I/O Read and Write (All)

Determines the time to read or write a block of data on the disk

## Event Dispatch (All)

Measures time to dispatch work between two running threads

## Float Rate (All)

Measures the rate of a mix of floating point operations

## Garbage Generation (Java)

Measures effectiveness and impact of GC on running threads





## Tests Continued

### Int Rate (All)

Measures the rate of a mix of integer instructions

### Lock Consistency (All)

Measures uncontended lock acquisition at various priorities

### Mcache perf (All)

Measures the performance of the memory cache

### Multi-CPU (C++, Java)

Creates objects and does work in a periodic scheduling setup

### Multi-Thread (All)

Throughput of matrix multiplication between 1 and NR\_CPUS

### Net I/O (All)

Measures time to read a block of data from a network connection

### NHRT No Heap Real Time (RTSJ)

Measures consistency of NHRT threads running on a loaded system



## More Tests

### Periodic Event (All)

Measures how consistently periodic events are dispatched

### Scoped Memory (RTSJ)

Measures the cost of accessing objects within various scopes

### Thread Priority (All)

Verifies that higher-priority threads preempt lower-priority threads

### Time Accuracy (All)

Evaluates the accuracy of the time subsystem by sleeping with the high-resolution timer and then comparing with the clock


### Timer (All)

Measures accuracy of the high-resolution timer by checking for the minimum time a thread can sleep and deterministically wake up



# Cross Language


## Determinism Guaranteed-RT





# Cross Platform


## Determinism Guaranteed-RT Intel v. AMD





## Loaded vs. Unloaded

### Determinism Loaded v. Unloaded





## Uses for RTMB

A solid, across-the-board real-time comparison (system, OS, language)  
Regressions

Daily test.kernel.org RTMB runs...

So you can brag about how awesome your box is

So you [your company] can sell more software

So you [your company] can sell more services

So you [your company] can sell more hardware

Just in case /dev/null is looking a little empty



## Legal Statement

This work represents the view of the author and does not necessarily represent the view of IBM.

IBM, IBM (logo), e-business (logo), pSeries, e (logo) server, and xSeries are trademarks or registered trademarks of International Business Machines Corporation in the United States and/or other countries.

Linux is a registered trademark of Linus Torvalds.

Other company, product, and service names may be trademarks or service marks of others.